

# THE PRODUCER - CONSUMER PROBLEM



**Course :** Operating System

**Class ID:** 161859

**Supervisor:** Dr. Đỗ Quốc Huy

# OUR TEAM

1. Trần Nguyễn Mỹ Anh - 20235474 - Leader
2. Nguyễn Thu Huyền - 20235507
3. Trần Lê Hạ Đan - 20235483
4. Nguyễn Khánh Ly - 20235600

# TABLE OF CONTENTS

1 Introduction

2 Problem Analysis

3 Synchronization Mechanism

4 Algorithm Design

5 Test Cases

6 Conclusion

01

# INTRODUCTION




# OVERVIEW

**Definition:** A classical synchronization problem involving multiple processes sharing a common resource.

## Components:


- **Producers:** Generate data items and place them in a shared buffer.
- **Consumers:** Remove and process items from the buffer.
- **Shared Buffer:** The critical resource where data is temporarily stored



# OVERVIEW

## AI/ML Context:

- **CPU (Producer):** Loads images, performs decompression, and data augmentation.
- **GPU (Consumer):** Retrieves batches for model training.
- **Challenge:** Poor coordination leads to the GPU remaining idle, wasting computational resources



# PROBLEM ANALYSIS

02



# PROBLEM STATEMENT

- **Key Components:** The system consists of Producers (responsible for generating data items) and Consumers (who process or use those items).
- **Communication Mechanism:** Interaction occurs through a shared buffer that temporarily stores items until they are consumed.
- **Bounded-Buffer Constraints:** In this variant, buffer usage is limited by its capacity.
  - **Overflow:** Producers must wait when the buffer is full.
  - **Underflow:** Consumers must wait when the buffer is empty.
- **Mutual Exclusion:** This is a critical requirement ensuring that only one process or thread modifies the buffer at a time to prevent data corruption

# POTENTIAL ISSUES

## Concurrency Issues:

- **Race Conditions:** Multiple threads modifying shared variables (counters/indices) simultaneously.
- **Data Inconsistency:** Items may be overwritten or lost.
- **Deadlock & Starvation:** Threads waiting indefinitely for conditions that never become true

## Performance issues:

- **CPU Utilization:** Inefficient techniques like Busy Waiting waste processing resources because threads continuously check conditions instead of blocking.
- **Synchronization Overhead:** Frequent context switches and unnecessary wake-ups can significantly degrade system performance.
- **Scalability:** Performance may decrease as the number of producers and consumers increases

# POTENTIAL ISSUES

**Problem Objectives:** A correct and optimized solution must satisfy the following criteria:

1. **Ensure Mutual Exclusion** during all buffer access operations.
2. **Prevent Buffer Overflow and Underflow** through proper coordination.
3. **Avoid Deadlock and Starvation** to ensure system progress.
4. **Optimize CPU Usage** by using blocking mechanisms instead of busy waiting.
5. **Support Scalability** by allowing multiple concurrent producers and consumers.

# SYNCHRONIZATION MECHANISMS



# SYNCHRONIZATION MECHANISMS

**Critical Section:** The code segment accessing the shared resource.

**Key Requirements:**

1. **Mutual Exclusion:** Only one thread executes a critical section at a time.
2. **Condition Synchronization:** Coordination based on buffer state (Full/Empty)

# THREE APPROACHES

Having established the core difficulties and the need for synchronization in the Producer-Consumer model, we now turn to evaluating three specific mechanisms: Busy Waiting, Blocking Synchronization (Condition Variables), and Semaphores.

1

Busy Waiting

2

Blocking  
Synchronization and  
Condition Variables

3


Semaphores

# ALGORITHM DESIGN

04




# BUSY WAITING



**Producer:** Prior to production, it continuously queries:  
“Is the buffer full?”

If it is full, the Producer waits and repeatedly re-queries until space becomes available.



**Consumer:** Prior to consumption, it continuously queries:  
“Is the buffer empty?”

If the buffer is empty, the Consumer waits and repeatedly re-queries until data is available.

# ALGORITHM

## Input Data:

- Shared Resource: A simple Python list buffer = [] acting as the storage.
- Constraints: BUFFER\_SIZE = 5
- Control flags: A global variable **count** is used to check the size to determine the current state of the Buffer.

To implement this method:

- Utilize an **empty while** loop to hold the process execution flow while the condition remains unsatisfied

---

### Algorithm 1: Busy Waiting

---

**Input :** Shared Buffer  $B$  with capacity  $N$

**Output :** Data transfer from Producer to Consumer

// Global variable: count = current number of items

```

1 Function Producer():
2   while True do
3     item ← produce_item();
4     // Busy Wait: Loop
5     // continuously while buffer is full
6     while count == N do
7       // Do nothing, just burn CPU cycles
8     B[count] ← item;
9     count ← count + 1;
10
11 Function Consumer():
12   while True do
13     // Busy Wait: Loop
14     // continuously while buffer is empty
15     while count == 0 do
16       // Do nothing
17     item ← B[count - 1];
18     count ← count - 1;
19     process_item(item);

```

---

# EVALUATION


- **Strengths**

- Simple and easy to implement
- No context switch overhead (can be fast in rare short-wait cases)

- **Weaknesses**

- Severe CPU waste due to constant polling
- Race conditions may cause data corruption
- Poor scalability as producers/consumers increase

# SLEEP & WAKEUP MECHANISM (CONDITION VARIABLES)



**Locking:** Before accessing the shared buffer, the thread must acquire a lock.

**Waiting:** If the condition is not met (e.g., buffer is full), the thread calls **wait()**. This action releases the lock and pauses the thread's execution.

**Signaling:** When a thread changes the buffer's state (e.g., adds an item), it calls **notify()** (or **signal()**) to wake up sleeping threads.

# ALGORITHM

## Input data:

- Synchronization Object: **threading.Condition()**. This object acts as both a Lock (Mutex) and a waiting room for threads.
- Shared Resource: The same **buffer** list with size 5.

## Program output:

- [Producer] Buffer Full. Waiting... (Thread stops here).
- [Consumer] Consumed item. Notifying...
- [Producer] Woke up. Produced item.

---

### Algorithm 2: Sleep & Wakeup (Condition Variables)

---

```
Input : Buffer B, Monitor Lock L,  
        Condition Variable CV

1 Function Producer():  
2     while True do  
3         item ← produce_item();  
4         AcquireLock(L);  
5         // Enter Critical Section  
6         while size(B) == N do  
7             // Release lock L and go to  
             sleep  
8             Wait(CV);  
9             push(B, item);  
10            // Wake up sleeping Consumer  
11            Notify(CV);  
12            ReleaseLock(L);

10 Function Consumer():  
11    while True do  
12        AcquireLock(L);  
13        // Enter Critical Section  
14        while size(B) == 0 do  
15            // Release lock L and go to  
            sleep  
16            Wait(CV);  
17            item ← pop(B);  
18            // Wake up sleeping Producer  
19            Notify(CV);  
20            ReleaseLock(L);  
21            process_item(item);
```

---

# ALGORITHM

- **Wait(condition):** This is the key difference from Method 1. Instead of looping, the thread stops execution here, saving CPU resources.
- **Signal(condition):** This command ensures that if the other party is sleeping, they will be alerted that the state has changed (e.g., from Empty to Not Empty).

---

**Algorithm 2:** Sleep & Wakeup  
(Condition Variables)

**Input :** Buffer  $B$ , Monitor Lock  $L$ ,  
Condition Variable  $CV$

```
1 Function Producer():
2   while True do
3     item ← produce_item();
4     AcquireLock(L);
5     // Enter Critical Section
6     while size(B) == N do
7       // Release lock L and go to
8       sleep
9       Wait(CV);
10    push(B, item);
11    // Wake up sleeping Consumer
12    Notify(CV);
13    ReleaseLock(L);

14 Function Consumer():
15   while True do
16     AcquireLock(L);
17     // Enter Critical Section
18     while size(B) == 0 do
19       // Release lock L and go to
20       sleep
21       Wait(CV);
22     item ← pop(B);
23     // Wake up sleeping Producer
24     Notify(CV);
25     ReleaseLock(L);
26     process_item(item);
```

---

# EVALUATION

- **Strengths**

- **CPU efficient:** no busy waiting while threads are blocked
- **Thread-safe:** mutex/locks prevent race conditions

- **Weaknesses**


- Context switch overhead when blocking/waking threads
- More complex logic; incorrect wait/notify may cause deadlocks
- Wake up without a signal → Require condition checks inside **while**,  
not **if**

# SEMAPHORE & MUTEX

**Idea:** An optimized solution focusing on managing permits to access resources

## Primitives Used:

- ***empty\_slots***: Counting semaphore initialized to Buffer Size
- ***full\_slots***: Counting semaphore initialized to 0
- ***mutex***: Binary semaphore/lock initialized to 1 for mutual exclusion



**Operations:** P (Wait) decrements the count/blocks; V (Signal) increments/wakes

# ALGORITHM

## Input data:

- **empty\_slots** : A *threading.Semaphore* initialized to BUFFER\_SIZE (5). Represents the number of free spaces.
- **full\_slots**: A *threading.Semaphore* initialized to 0. Represents the number of items ready for consumption.
- **mutex**: A *threading.Lock()* to protect the data integrity of the list.

## Program output:

If the Producer is faster, the logs will show [Producer] Waiting for empty slot... and it will strictly wait until a [Consumer] Consumed... log appears

**Algorithm 3:** Semaphores  
(Optimized Solution)

```
Input : Semaphores: Empty (init N),  
        Full (init 0), Mutex (init 1)  
1 Function Producer():  
2 while True do  
3     item  $\leftarrow$  produce_item();  
        // Step 1: Wait for an empty  
        slot  
4     Wait(Empty);  
        // Step 2: Enter Critical  
        Section (Exclusive Access)  
5     Wait(Mutex);  
6     add_to_buffer(item);  
7     Signal(Mutex);  
        // Step 3: Signal that a new  
        item is available  
8     Signal(Full);  
9 Function Consumer():  
10    while True do  
        // Step 1: Wait for available  
        data  
11    Wait(Full);  
        // Step 2: Enter Critical  
        Section  
12    Wait(Mutex);  
13    item  $\leftarrow$  remove_from_buffer();  
14    Signal(Mutex);  
        // Step 3: Signal that a slot is  
        now empty  
15    Signal(Empty);  
16    process_item(item);
```

# ALGORITHM

- **Separation of Concerns:** The **Wait(empty\_slots)** and **Signal(full\_slots)** pair in the Producer manages the flow control, while **Wait(mutex)** handles data integrity.
- **No spurious wakeups:** Unlike Method 2, Semaphores generally do not suffer from spurious wakeups, making the logic cleaner (no While loop needed for the lock itself, though the semaphore handles the blocking).

**Algorithm 3: Semaphores (Optimized Solution)**

**Input :** Semaphores: *Empty* (init *N*), *Full* (init 0), *Mutex* (init 1)

```
1 Function Producer():
2   while True do
3     item ← produce_item();
4     // Step 1: Wait for an empty slot
5     Wait(Empty);
6     // Step 2: Enter Critical Section (Exclusive Access)
7     Wait(Mutex);
8     add_to_buffer(item);
9     Signal(Mutex);
10    // Step 3: Signal that a new item is available
11    Signal(Full);

12 Function Consumer():
13   while True do
14     // Step 1: Wait for available data
15     Wait(Full);
16     // Step 2: Enter Critical Section
17     Wait(Mutex);
18     item ← remove_from_buffer();
19     Signal(Mutex);
20     // Step 3: Signal that a slot is now empty
21     Signal(Empty);
22     process_item(item);
```

# EVALUATION

- **Strengths**

- Highly scalable: supports multiple Producers and Consumers efficiently
- Deadlock-safe with correct semaphore ordering
- Industry-standard model used in message queues and AI data pipelines

- **Weaknesses**

- Higher implementation complexity
- Semaphore misordering may cause deadlock
- Debugging synchronization errors is difficult

# TEST CASES



## SCENARIO 1: BUFFER OVERFLOW

- **Configuration:**
  - **Buffer size** = 2.
  - **Producer sleep time** = 0.1s.
  - **Consumer sleep time** = 1.0s.
- **Expected behavior:** The Producer should fill the buffer quickly and then enter a “Waiting” state. It must not overwrite existing data.

```
--- DEMO METHOD 1: BUSY WAITING & HIGH CPU ---
```

```
Program will run in 15 seconds...
```

```
[MONITOR] Start CPU Monitoring... (Wait 1-2 seconds for stable readings)
[SYSTEM WARNING] High CPU Usage: 69.2% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.5% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.4% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 86.1% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.3% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.5% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 83.6% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.8% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.5% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.6% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.6% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.8% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 89.6% (Busy Waiting detected!)
[SYSTEM WARNING] High CPU Usage: 90.1% (Busy Waiting detected!)
```

```
--- STOP PROGRAM ---
```

```
[SYSTEM WARNING] High CPU Usage: 74.8% (Busy Waiting detected!)
```

## CPU usage - Method 1

**Busy Waiting (Method 1):** CPU usage spiked to 90.1% as the Producer loop checked the condition continuously

## CPU usage - Method 3

### Semaphores :

The Producer paused  
efficiently.

```
--- DEMO METHOD 3: SEMAPHORES & LOW CPU ---
Program will run for 15 seconds...

[MONITOR] Start CPU Monitoring... (Wait 1-2 seconds for stable readings)
[SYSTEM] Excellent! CPU Usage: 0.1% (Optimized)
[SYSTEM] Excellent! CPU Usage: 0.0% (Optimized)
[SYSTEM] Excellent! CPU Usage: 0.3% (Optimized)
[SYSTEM] Excellent! CPU Usage: 0.0% (Optimized)
[SYSTEM] Excellent! CPU Usage: 0.2% (Optimized)

--- STOP PROGRAM ---
[SYSTEM] Excellent! CPU Usage: 0.1% (Optimized)
```

Console log showed: [Producer] Waiting for empty slot....

## SCENARIO 2: BUFFER UNDERFLOW

- **Configuration:**
  - Buffer size N = 5.
  - Producer sleep time = 1.0s.
  - Consumer sleep time = 0.1s.
- **Expected behavior:** The Consumer should empty the buffer and then wait. It must not crash or retrieve None/garbage data.

# SCENARIO 2: BUFFER UNDERFLOW

```
--- TEST SCENARIO 2: BUFFER UNDERFLOW (CONSUMER > PRODUCER) ---
Config: Buffer=5, Producer Sleep=1.0s, Consumer Sleep=0.1s

[13:28:20] [Producer] Added 42.      Buffer: [42] (Len: 1)
[13:28:20] [Consumer] Waiting for data...
[13:28:20] [Consumer] Consumed 42.    Buffer: [] (Len: 0)
[13:28:20] [Consumer] Waiting for data...
[13:28:21] [Producer] Added 45.      Buffer: [45] (Len: 1)
[13:28:21] [Consumer] Consumed 45.    Buffer: [] (Len: 0)
[13:28:21] [Consumer] Waiting for data...
[13:28:22] [Producer] Added 51.      Buffer: [51] (Len: 1)
[13:28:22] [Consumer] Consumed 51.    Buffer: [] (Len: 0)
[13:28:22] [Consumer] Waiting for data...
[13:28:23] [Producer] Added 58.      Buffer: [58] (Len: 1)
[13:28:23] [Consumer] Consumed 58.    Buffer: [] (Len: 0)
[13:28:23] [Consumer] Waiting for data...
[13:28:24] [Producer] Added 28.      Buffer: [28] (Len: 1)
[13:28:24] [Consumer] Consumed 28.    Buffer: [] (Len: 0)
[13:28:24] [Consumer] Waiting for data...
[13:28:25] [Producer] Added 17.      Buffer: [17] (Len: 1)
[13:28:25] [Consumer] Consumed 17.    Buffer: [] (Len: 0)
[13:28:25] [Consumer] Waiting for data...
[13:28:26] [Producer] Added 70.      Buffer: [70] (Len: 1)
[13:28:26] [Consumer] Consumed 70.    Buffer: [] (Len: 0)
[13:28:26] [Consumer] Waiting for data...
[13:28:27] [Producer] Added 100.     Buffer: [100] (Len: 1)
[13:28:27] [Consumer] Consumed 100.   Buffer: [] (Len: 0)
[13:28:27] [Consumer] Waiting for data...
[13:28:28] [Producer] Added 5.       Buffer: [5] (Len: 1)
[13:28:28] [Consumer] Consumed 5.    Buffer: [] (Len: 0)
[13:28:28] [Consumer] Waiting for data...
[13:28:29] [Producer] Added 71.      Buffer: [71] (Len: 1)
[13:28:29] [Consumer] Consumed 71.    Buffer: [] (Len: 0)
[13:28:29] [Consumer] Waiting for data...

--- STOPPING TEST ---
```

**Result :** The Consumer paused correctly.

Console log showed: *[Consumer] Waiting for data....*

## SCENARIO 3: CONCURRENCY & DATA INTEGRITY

- **Configuration:** 5 Producers and 5 Consumers operating simultaneously.
- **Expected behavior:** All items produced must be consumed exactly once. No items should be lost or duplicated.

```
---- TIME'S UP! STOPPING PRODUCERS... ----  
-----  
FINAL RESULT (DATA INTEGRITY CHECK):  
Total Items Produced: 8992  
Total Items Consumed: 8992  
  
>>> SUCCESS: PERFECT MATCH! NO DATA LOSS/DUPLICATION. <<<
```

### Result:

Using method 3 (Semaphores + Mutex), the total count of produced items matched the total count of consumed items perfectly after a 60-second run

- **The most robust and suitable technique for real-world systems, offering the best balance of performance and scalability**

THANK  
YOU

DO YOU HAVE ANY  
QUESTIONS?